


- How does the computer fetch an instruction?

- 1. Start processing instruction
- 2. Generate data address
- 3. Compute page number
- 4. If page is in memory
 - then
 - get data and finish instruction
 - advance to the next instruction
 - return to step 1
 - else
 - generate page interrupt
 - call page interrupt handler

• Algorithm: Page Interrupt Handler

- 1. If there is no free page frame
- then
- select page to be swapped out using
- a page removal algorithm
- update job's Page Map Table → 1st update
- if content of page had been changed } lazy write
- then write page to disk } is also possible.
- 2. Use page number (step 3 of the previous
- algorithm to get disk address where
- page is stored (the File Manager, to
- be discussed later, uses the page
- number to get the disk address)
- 3. Read page into memory
- 4. Update job's Page Map Table ← 2nd update, for the incoming page's j=b.
- 5. Update Memory Map Table
- 6. Restart interrupted instruction

- Although demand paging is a solution to inefficient memory utilization, it does not solve all the problems
- **Thrashing:** if a large amount of page swapping is performed, the system efficiency is affected.
- **Page fault:** a failure to find a page in memory.



CSCI 460—Operating Systems

Lecture 4

Memory Management—memory hierarchy

Textbook: Operating Systems
by William Stallings

1. The Memory Hierarchy

- In the more recent decades, computer memory is not arranged in a linear fashion.
- The design constraints on memory rest on:

- 1. Capacity.
- 2. Speed (access time).
- 3. Cost (unit cost).

- Their relationship

- Faster Speed (access time) → Greater Cost.
- Greater Capacity → Smaller Cost.
- From these two, we have: Greater Capacity → Slower Speed.
- So you can't have Greater Capacity, Small Cost and Fast Speed at the same time!

$A \rightarrow B \Leftrightarrow \sim B \rightarrow \sim A$
or, Smaller cost → slower speed

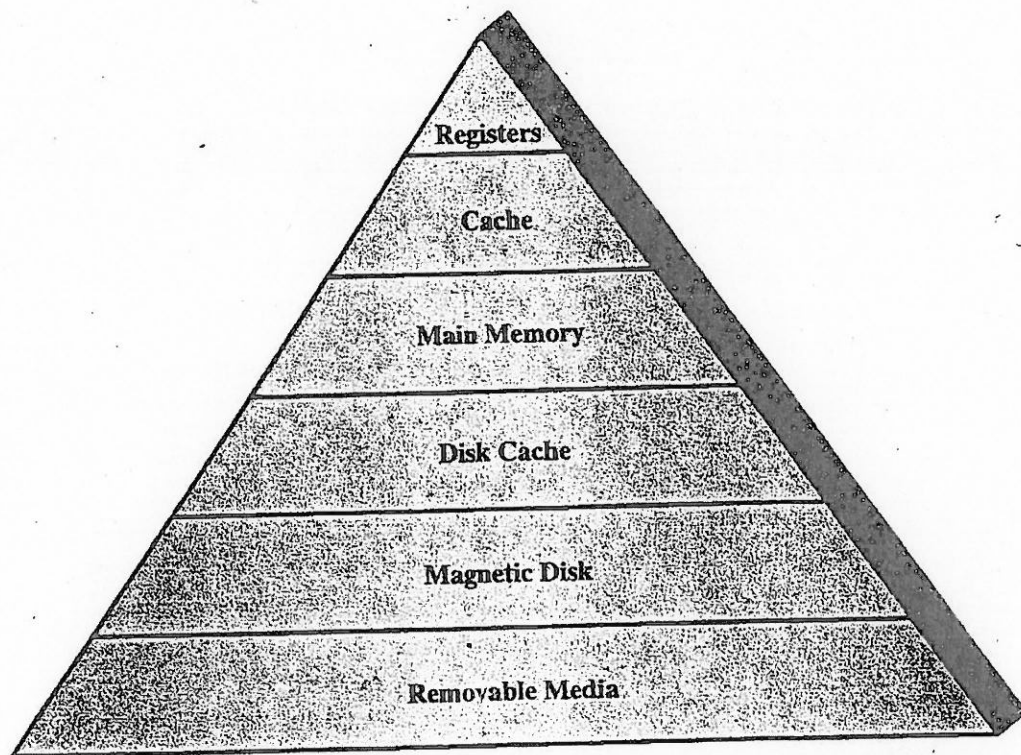
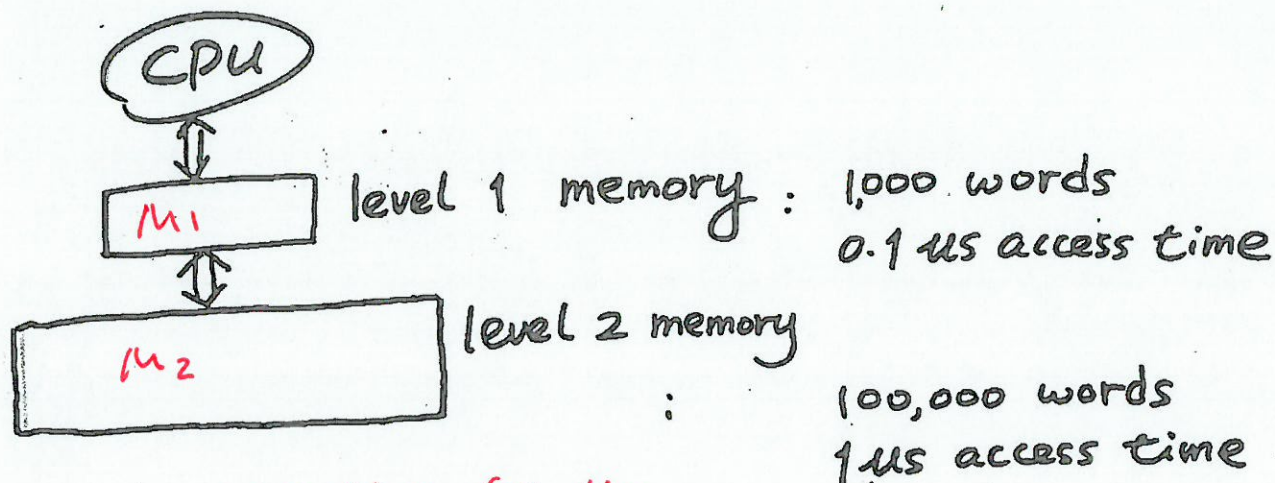


Figure 1.14 The Memory Hierarchy

2. Memory Hierarchy (cont.)

- If we look from top to bottom at Figure 1.14 (in Stallings), the following can be observed.
 - Cost is decreasing.
 - Memory capacity is increasing.
 - Speed is slowing down.
 - Frequency of access of memory by processor is decreasing.
- Why?
 - **Locality of Reference.**
 - Locality of reference is not only valid in OS. It is the basis for compiler optimization, computer architecture and database management (and recently in the Internet browsing).
- Thanks to the semiconductor industry (for building different kinds of storage media for us)!



T_1 — access time for M_1

T_2 — access time for M_2

T — access time for the whole system

H — Hit ratio, the probability that you can find an item in M_1 .

$$T = H \cdot T_1 + (1-H) \cdot (T_1 + T_2)$$

$$= \cancel{H T_1} + T_1 + T_2 - \cancel{H T_1} - H T_2$$

$$= T_1 + T_2 - H T_2$$

$$= T_1 + (1-H) \cdot T_2$$

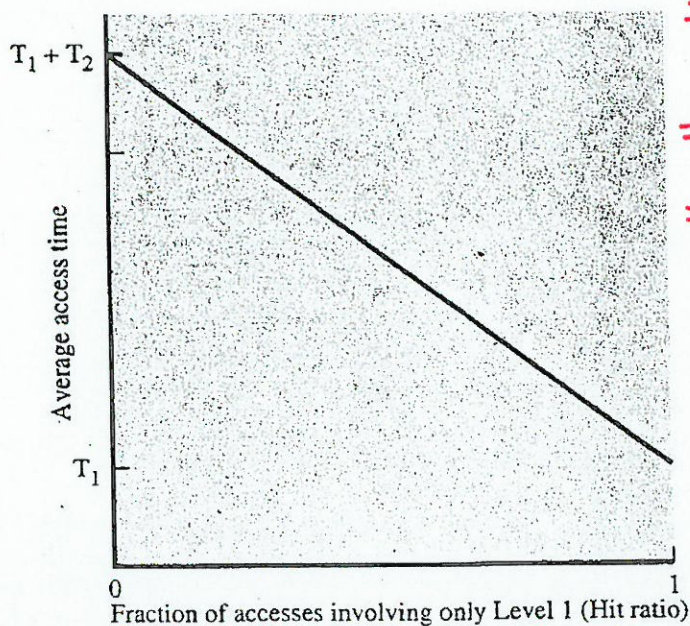


Figure 1.15 Performance of a Simple Two-Level Memory


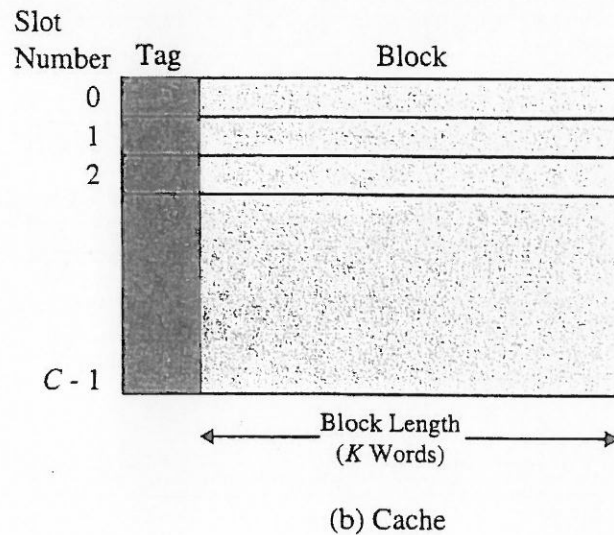



Figure 1.16 Cache and Main Memory



$$C \ll M = \frac{2^n}{K}$$

of blocks
in main memory

Figure 1.17 Cache/Main-Memory Structure

3. Cache Memory

- **Motivation.**

- On all instruction cycles, the processor access memory at least once: to fetch the instruction, to fetch operands and/or store the results. *Think of executing an assemble instruction: ADD C, A, B ($C \leftarrow A + B$).*
 - In general memory access speed cannot match the processor speed. So it makes sense to exploit the principle of locality by building a small, fast memory between the processor and main memory.
- This fast memory, almost *invisible* from OS, is **cache**.
 - The objective of cache memory is to speed up the memory so that it is almost as fast as the speed of processor and at the same time it provides a memory size which is large enough (for most jobs).
 - Let's us look at the structure of a cache/memory system.

RA — Read Address

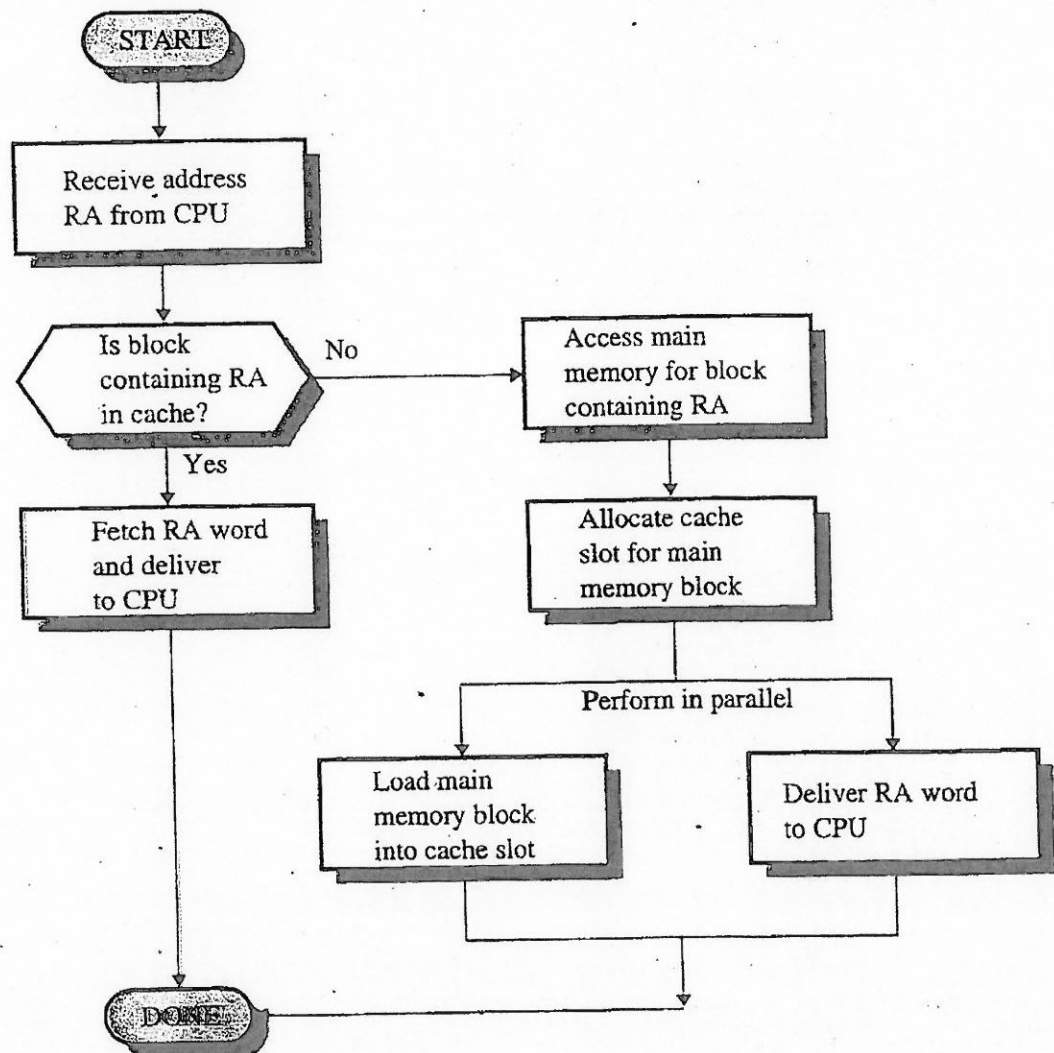


Figure 1.18 Cache Read Operations

4. Cache Memory (cont.)

- Let's look at Figure 1.18. What problems can you see with this example?

*- Sth must be moved out of the cache to load the block containing R/A into the cache.
- write/wait?*


- Cache design is beyond this course. But the following issues must be considered in general.

- 1. Cache size.
- 2. Block size. Suitable size of block will ensure that the hit ratio is high.
- 3. Mapping function. When a block is read into the cache, the 1st question is to decide where we should put it. (2 hints: (A) When one block is read in, another one should be moved out, so we should minimize the probability that a moved-out block will be referenced again in the near future. (B) The more flexible the mapping function, the more time it takes to search the cache to find a block.)
- 4. The replacement policy. (Can you think of some?)
- 5. Write policy. If the contents of a block in the cache are changed, we should write it back to the main memory before replacing it. So when should this write operation takes place?

hash function without chaining

FIFO

5 Performance Analysis of Two-level Memory



- Assume that we have two levels of memory, M_1, M_2 (M_1 is smaller, but faster.) Let's first look at the average system access time T_s .

$$\begin{aligned} T_s &= H \times T_1 + (1 - H) \times (T_1 + T_2) \\ &= T_1 + (1 - H) \times T_2 \end{aligned} \quad (1.1)$$

where

T_s = average (system) access time

T_1 = access time of M_1 (e.g., cache, disk cache)

T_2 = access time of M_2 (e.g., main memory, disk)

H = hit ratio (fraction of time reference is found in M_1)

- Let $\frac{T_1}{T_s}$ be the *access efficiency*, we have

$$\frac{T_1}{T_s} = \frac{1}{1 + (1 - H) \frac{T_2}{T_1}} \rightarrow \text{Constant.}$$

We want this ratio to be close to 1.




Figure 1.23 Access Efficiency as a Function of Hit Ratio ($r = T_2/T_1$)

- Let's now look at the average cost per bit for the two-level memory, C_s .

$$C_s = \frac{C_1 S_1 + C_2 S_2}{S_1 + S_2} \quad (1.2)$$

where

C_s = average cost per bit for the combined two-level memory

C_1 = average cost per bit of upper-level memory M1

C_2 = average cost per bit of lower-level memory M2

S_1 = size of M1

S_2 = size of M2

- To make C_s roughly the same as C_2 . We should make $S_1 \ll S_2$. ($C_1 \gg C_2$ due to the hardware cost, which we can do very little to change it.) Notice that \rightarrow constant

$$\frac{C_s}{C_2} = \frac{\frac{C_1}{C_2} + \frac{S_2}{S_1}}{1 + \frac{S_2}{S_1}}$$




Figure 1.22 Relationship of Average Memory Cost to Relative Memory Size for a Two-Level Memory

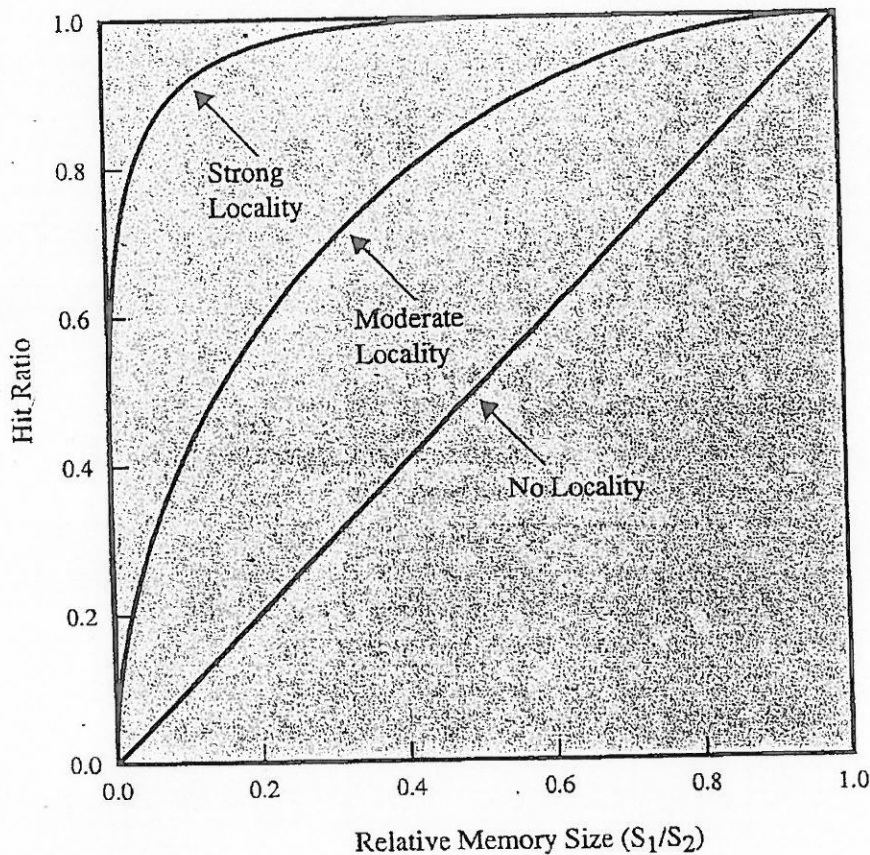


Figure 1.24 Hit Ratio as a Function of Relative Memory Size

In practice,

- ① Cache size: 1K ~ 128K
- ② Hit ratio > 0.75 almost all the time

